

Contents

LostSideDead ACMX2 Explained	2
acidcam-gpu + ACMX2 Complete Technical Reference	2
Program Screenshot	2
Project Origin and Purpose	2
Why CUDA Kernels and GLSL Shaders?	3
How Each Technology Is Used	3
ACMX2 Qt Application Orchestration Layer	4
How the Qt Interface Works	4
Standalone CLI acidcam-gpu/app/main_cv.cu	4
Startup: Device and Camera Discovery	5
AnimationState The Per-Frame Parameter Evolution Engine	5
updateAndDraw() The Frame Processing Function	5
Output Encoding via MXWrite	5
The Public API ac-gpu.hpp	6
Key Types and Definitions	6
The DynamicFrameBuffer Class	6
The launch_filter() Function	6
FilterParams The Per-Frame Parameter Bundle	7
The Unified Kernel unifiedFilterKernel	8
Thread Layout and Pixel Assignment	8
The Filter Chain Loop	8
Device Helper Functions	8
The Filter Catalog 736 Effects Across All Categories	9
Temporal / Trail Effects	9
Geometric Distortion	9
Color Manipulation	9
Bitwise / XOR Operations	10
Glitch and Noise	10
Scan / Line Effects	10
Block / Square Effects	10
Wave / Oscillation Effects	11
Mirror / Reflection	11
Pixel Read / Strobe / Blend Collections	11
Advanced / Cinematic Effects	11
The Two-Stage GPU Pipeline CUDA Compute + GLSL Shaders	12
Stage 1: CUDA Compute Kernels Massively Parallel Pixel Processing	12
Stage 2: GLSL Shaders Full-Frame Post-Processing	12
The Transfer: Zero-Copy GPU Interop	13
Filter and Shader Stacking Two Independent Composable Chains	13
CUDA Filter Chain (Stacking Inside the Kernel)	13
GLSL Shader Pass Stack (Ping-Pong Framebuffer Passes)	13
Combined Stacking: The Full Pipeline	14
Why Stacking Produces Infinite Visuals	14
In-Depth: Why Order Dominates Output	14
Build System, Dependencies, and Distribution	14
Required Dependencies	15
Build Order	15
Container Distribution (Podman)	15
Development Environment	15
Source Code Map by Project Part	15
ACMX2 Core	15

ACMX2 Media + Audio + 3D	16
ACMX2 Interface + Tools	16
acidcam-gpu Library	16
Ops + Deployment	16
Expanded Meaning of Each Code Map Item	16
End-to-End Runtime Flow	17
Detailed Runtime Inner Workings	17
Project Architecture Goals and Design Tradeoffs	17
1) Composability First	17
2) Deterministic Runtime Paths	17
3) Throughput-Oriented Execution	18
4) Operator Visibility	18
Engineering Implications	18
Build, Deployment, and Operational Model	18
Build Surface	18
Operational Surface	18
Project Summary	18

LostSideDead ACMX2 Explained

Original page: <https://lostsidedead.biz/acmx2-explained.html>

acidcam-gpu + ACMX2 Complete Technical Reference

acidcam-gpu is a CUDA-accelerated real-time video effects engine. It contains a library of **736 GPU filter kernels** that run entirely on the NVIDIA GPU, a **unified dispatch kernel** that chains any ordered subset of those filters per frame, a rotating device-side frame history buffer for temporal effects, and a CLI application that drives the full pipeline from camera or file input through to live display and optional MXWrite output encoding.

ACMX2 is the Qt-based orchestration layer it wraps the same library and CLI with a visual interface, shader pass ordering, session persistence, and process monitoring. Together they form a layered creative tool: CUDA for pixel-level power, OpenGL for display, Qt for control, and Podman containers for distribution.

The system operates as a **two-stage GPU pipeline**: first, CUDA kernels process every pixel in parallel on the GPUs streaming multiprocessors (one thread per pixel), applying a user-defined chain of filters entirely in device memory. Then, the CUDA output is transferred to OpenGL via a Pixel Buffer Object (PBO) a zero-copy GPU-to-GPU transfer where it becomes a texture. GLSL fragment shaders then apply additional full-frame visual effects in one or more stacked passes using ping-pong framebuffers. Both the CUDA filter chain and the GLSL shader pass stack are independently configurable and orderable, producing a combinatorial explosion of possible visual outcomes.

Source: <https://github.com/lostjared/acidcam-gpu> C++20 / CUDA 12.x / OpenCV CUDA / OpenGL / Qt6 / MXWrite.

Program Screenshot

This is the current desktop ACMX2/acidcam-gpu interface in action.

Project Origin and Purpose

The Acid Cam project originated as a CPU-based filter library (**libacidcam**) that applied artistic glitch effects to camera and video frames. As resolutions grew and filter stacks became more complex, CPU-based pixel processing became a bottleneck applying dozens of per-pixel transforms on high-resolution video ate into frame budget quickly.




Figure 1: ACMX2 and acidcam-gpu running on desktop

acidcam-gpu was created to solve this. The solution was to port the entire filter library to NVIDIA CUDA, so all pixel work runs massively in parallel on the GPU. A unified dispatch kernel allows any ordered combination of the 736 available filters to run in a single GPU pass per frame, without recompiling. The result is a system capable of applying complex multi-layer visual transforms at full framerate on modern NVIDIA hardware.

Why CUDA Kernels and GLSL Shaders?

Both CUDA kernels and GLSL shaders execute on the **same physical hardware** the GPUs streaming multiprocessors (often marketed as CUDA cores). The difference is the **programming model**, not the silicon. CUDA is NVIDIA's general-purpose compute API: kernels can freely read from arbitrary device memory, perform complex conditional logic, access a ring buffer of historical frames, and do per-pixel computation without any graphics pipeline constraints. GLSL runs through the OpenGL graphics pipeline as fragment shader programs, which excel at texture sampling, interpolation, and full-frame image processing with ShaderToy-compatible uniforms like time, resolution, and mouse position. By using both programming models on the same GPU hardware, the system gets the best of each approach.

How Each Technology Is Used

- **CUDA (pixel-level compute):** The 736 filter kernels run as CUDA threads one thread per pixel across the GPUs streaming multiprocessors simultaneously. A 1920x1080 frame launches over 2 million threads in parallel. The unified kernel loops through the user's selected filter chain, applying each filter sequentially to the pixel data in-place. Filters can read from a ring buffer of prior frames stored entirely in GPU memory, enabling temporal effects like **MedianBlend** (averaging across all history frames), **AuraTrails** (blending against frames at indices 1, 4, and 7), and **MatrixOutline** (comparing against a frame from 4 steps back). This kind of arbitrary multi-frame memory access is natural in the CUDA compute model but difficult or impossible within the GLSL fragment shader pipeline, which is designed around single-texture-per-pass processing.
- **GLSL (full-frame shader effects):** After CUDA processing, the result is transferred to an OpenGL

texture via PBO interop (no CPU round-trip). GLSL fragment shaders running on the same GPU hardware via the OpenGL graphics pipeline then process the entire frame as a texture, applying effects like color grading, distortion, glow, CRT simulation, and other post-processing transforms. The shader library supports Shadertoy-compatible uniforms (`iTime`, `iResolution`, `iMouse`, `iFrame`, etc.) and can render onto 2D quads or 3D model geometry via MXMOD meshes.

- **Composition:** CUDA filters and GLSL shaders compose at the frame level. CUDA runs first on the raw pixel data, then GLSL operates on the CUDA-filtered result. Both stages support stacking multiple CUDA filters chain inside a single kernel launch, and multiple GLSL shaders chain via ping-pong framebuffer passes. The full pipeline is: Camera CUDA filter chain PBO transfer GLSL shader pass 1 GLSL shader pass 2 Screen.

ACMX2 Qt Application Orchestration Layer

ACMX2 is the Qt6 application that wraps the `acidcam-gpu` CLI and library into a controllable desktop session. It handles shader pass management, GPU filter ordering, session persistence, and process lifecycle it does not re-implement any pixel processing itself.

- **Main window:** `ACMX2/interface/main_window.cpp` QProcess supervision, log rendering, session menu
- **GPU filter dialog:** `ACMX2/interface/gpufilter.cpp` calls `--list-filters`, parses output, stores ordered selection
- **Shader pass ordering:** `ACMX2/interface/shaderpass.cpp` explicit multi-pass chain, output of pass N feeds into pass N+1
- **Audio integration:** `ACMX2/audio.cpp/.hpp` RtAudio amplitude extraction, reactive parameter modulation
- **Shader cache:** `ACMX2/program.cpp/.hpp` binary cache with source+driver fingerprinting to skip recompile
- **3D geometry:** `ACMX2/models/*.mxmod` MXMOD-format geometry used in scene-influenced render stages
- **Build:** `ACMX2/CMakeLists.txt` and `ACMX2/interface/CMakeLists.txt` validates CUDA, OpenCV CUDA headers, FFmpeg, MXWrite, Qt6, RtAudio at configure time

How the Qt Interface Works

1. **Filter discovery:** `gpufilter.cpp` spawns the `acmx2` binary with `--list-filters`, reads `index:name` lines from stdout, sorts the list alphabetically, and populates a selection dialog. The user picks from the full 736-filter catalog and reorders them via drag or list controls.
2. **Chain persistence:** selected filter order is stored as an ordered index list and passed back to `acmx2` as arguments at run time. Changing the order produces a completely different visual output without touching any code.
3. **Process supervision:** `main_window.cpp` binds a `QProcess` to the `acmx2` executable (built from `ACMX2/acmx.cpp`). stdout/stderr streams are captured and shown in the UI log panel so runtime errors (bad camera index, missing CUDA device, encoder failure) are immediately visible.
4. **Session settings:** `QSettings` persists executable path, shader directory, preferred styles, and last-used filter chain so sessions can be resumed exactly.
5. **Shader pass layer:** `shaderpass.cpp` manages an ordered list of GLSL pass configs. Each pass uses the previous pass output as input, compositing GLSL effects on top of the CUDA-processed frame.
6. **Audio-reactive path:** when RtAudio is enabled at compile time, an audio callback computes per-buffer amplitude average. This value can modulate alpha or other per-frame parameters, making output visually reactive to microphone or line input.

Standalone CLI `acidcam-gpu/app/main_cv.cu`

The standalone CLI app (`app/main_cv.cu`) is an independent tool that drives the CUDA filter pipeline directly from the command line without ACMX2. It handles input device management, argument parsing,

frame loop control, animation state evolution, and output encoding. It is written in CUDA C++ (compiled by nvcc) and targets C++20. Note: the ACMX2 Qt interface does **not** supervise this file instead, `QProcess` launches and monitors the `acmx2` executable (built from `ACMX2/acmx.cpp`), which contains its own runtime loop and GL/CUDA pipeline.

Startup: Device and Camera Discovery

- **checkDevices():** calls `cudaGetDeviceCount()`. If the count is zero or the call errors, it prints a detailed message (driver not installed, GPU not seated) and exits immediately. On success, prints the CUDA device short info via OpenCV's `cv::cuda::printShortCudaDeviceInfo()`.
- **listGraphicsCards():** enumerates all CUDA devices using `cv::cuda::DeviceInfo`, prints each device's name and total VRAM in MB.
- **listCameras():** probes `/dev/video0` through `/dev/video9` with `cv::VideoCapture`. During probing, `stderr` is redirected to `/dev/null` to suppress OpenCV's verbose device error messages. For each working device, reads its human-readable name from `/sys/class/video4linux/videoN/name`.
- **Signal handling:** a custom `Interrupt` exception class and a `signalHandler` that throws it allows `SIGINT/SIGTERM` to cleanly unwind the frame loop without leaving the CUDA device in a dirty state.

AnimationState The Per-Frame Parameter Evolution Engine

The `AnimationState` struct (global `gState`) drives the continuous evolution of visual parameters each frame. This is what makes the visuals non-static the same filter chain produces different pixel values every frame because the parameters feeding it change continuously:

- **alpha oscillation:** `alpha` starts at 1.0 and increments by 0.01f per frame until it reaches 3.0f, then decrements back to 1.0f. This creates a smooth breathing rhythm in blend-intensity across all alpha-parameterized filters. The full cycle takes 400 frames at 60fps 6.7 seconds.
- **frame index oscillation:** `current_frame_index` bounces between 0 and `arraySize - 1`, incrementing or decrementing by 1 each frame based on `index_dir`. This continuously sweeps which historical frame is selected by the `start_index` parameter effects that read from frame history change which past frame they reference every tick.
- **square_size oscillation:** `square_size` oscillates between 2 and 64 pixels, changing by 2 per frame. This controls block and tile dimensions for square-based effects. The oscillation causes visible growth and shrinkage in any block-decomposition filter.

updateAndDraw() The Frame Processing Function

Called every frame of the main loop, this function performs the full GPU pipeline:

1. Updates all three `AnimationState` values (`alpha`, `frame_index`, `square_size`) based on their current direction and bounds.
2. Copies device frame pointers from the `DynamicFrameBuffer`'s `rawPointers` vector into the device-side pointer array (`d_ptrList`) via `cudaMemcpy(HostToDevice)`.
3. Copies the most recently uploaded frame from the buffer into the working GPU buffer using `cudaMemcpy2D(DeviceToDevice)` this is a pure device-to-device copy, no round-trip through CPU RAM.
4. Calls `launch_filter()` with the current animation state, the working buffer, and the active filter list. The CUDA kernel runs and transforms the working buffer in-place.
5. The processed frame in the GPU working buffer is then downloaded or blitted for display/encoding.

Output Encoding via MXWrite

When recording is enabled, processed frames are passed to the `MXWrite` library which wraps FFmpeg encoding. `MXWrite` accepts raw frame data and handles container muxing, codec selection, and file writing. This keeps encoding off the critical path the CUDA filter loop is not blocked by disk I/O.

The Public API `ac-gpu.hpp`

The header `acidcam-gpu/include/ac-gpu/ac-gpu.hpp` defines the complete ABI contract between the host application and the CUDA engine. Everything the caller needs to know lives here.

Key Types and Definitions

- **AC_FILTER_MAX = 736** the total number of available filters. Filter indices run from 0 to 735 inclusive. This constant lets callers validate index bounds before dispatch.
- **struct GPUFilter { int index; }** the lightweight device-side representation of a filter. Only the integer index crosses to the GPU; names and metadata stay on the host.
- **struct Filter { int index; std::string name; GPUFilter toGPU() const; }** the host-side filter descriptor. `toGPU()` produces the compact `GPUFilter` for device transfer.
- **class ACException** thin exception type for runtime errors (bad resolution strings, missing devices). Carries a message string via `why()`.
- **CHECK_CUDA(call)** macro that wraps any CUDA runtime call, checks the error code, prints file/line/message, and exits on failure. Used throughout both the library and the CLI app for consistent error handling.
- **extern Filter filters[]** the master filter table defined in `filters.cu`, exposed for host-side iteration. Callers can walk this array to build filter lists, populate UI controls, or serialize chain configs.

The `DynamicFrameBuffer` Class

`DynamicFrameBuffer` manages a ring of historical frames entirely in device memory (`cv::cuda::GpuMat`). This is the mechanism that makes temporal effects possible filters can read from any prior frame in the ring without downloading anything to host memory.

- **Construction:** takes an `arraySize` parameter (e.g., 8) that sets ring depth. Internally creates a `std::vector<cv::cuda::GpuMat>` and a parallel `std::vector<unsigned char*>` of raw device pointers.
- **update(const cv::Mat& inputFrame):** 1. Uploads the CPU frame into `d_uploadBuffer` (a staging `GpuMat`) via `d_uploadBuffer.upload(inputFrame)`. 2. On resolution change, reallocates all ring slots as `CV_8UC4` (8-bit per channel, 4 channels RGBA) at the new dimensions and resets `completedFrames = 0`. 3. Calls `std::rotate(deviceFrames.begin(), deviceFrames.begin()+1, deviceFrames.end())` this shifts the oldest frame to position `back()` without copying pixel data, just rotating smart handles. 4. Converts and copies the new frame into the back slot: if the upload buffer is 3-channel BGR (standard OpenCV camera output), `cv::cuda::cvtColor(d_uploadBuffer, deviceFrames.back(), cv::COLOR_BGR2RGBA)` converts to 4-channel RGBA on the GPU. If already 4-channel, a direct `copyTo` is used. 5. Syncs `rawPointers[i] = deviceFrames[i].data` for all slots, giving the kernel a plain raw pointer into each frame's device memory. 6. Increments `completedFrames` up to `arraySize`, tracking how many valid (non-zero) history entries exist before the ring fills.
- **getDeviceFramePointers():** returns the `rawPointers.data()` a host-side array of device pointers. This is copied to a device-side pointer array before kernel launch so the kernel can dereference historical frame data.
- **Memory layout:** each `cv::cuda::GpuMat` is a pitched allocation. The `framePitch` value (row stride in bytes) is stored on first allocation and passed to kernel as `step`. Pitched memory improves coalescing for row-based access patterns.

The `launch_filter()` Function

This is the C-linkage function that the host calls every frame to execute the filter chain. It handles lazy filter list management and kernel dispatch:

1. **Guard:** returns immediately if `c == 0` (no filters), dimensions are zero, or `numFrames` is zero.

2. **Lazy filter list rebuild:** if `changed == true` or the device filter list pointer is null, it synchronizes the device (`cudaDeviceSynchronize`), frees any existing device list, converts host `Filter[] GPUFilter[]` using `toGPU()`, allocates a new device array with `cudaMalloc`, and copies with `cudaMemcpy(HostToDevice)`. Sets `changed = false` after update. This means re-ordering the filter chain is essentially free at runtime the rebuild only happens when the chain actually changes.
3. **Grid/block computation:** `dim3 blockSize(16, 16)` 256 threads per block arranged in a 2D tile. `dim3 gridSize((width+15)/16, (height+15)/16)` enough tiles to cover the full frame, rounded up so edge pixels are handled correctly.
4. **Per-frame parameter assembly:** a fresh `FilterParams` struct is built each frame using the current animation state and fresh random values from `rand()`. This per-frame randomness is what makes the visuals continuously evolve without any explicit animation scripting.
5. **Kernel dispatch:** `unifiedFilterKernel<<<gridSize, blockSize>>>(d_list, c, data, allFrames, w, h, step, params)`.
6. **Synchronization:** `cudaDeviceSynchronize()` is called after the kernel launch to ensure the frame is fully processed before the result is read or displayed.

FilterParams The Per-Frame Parameter Bundle

The `FilterParams` struct is defined inside the `ac_gpu` namespace in `filters.cu` and is passed by value into the kernel. Every filter case inside the kernel reads from this shared param bundle rather than maintaining its own state, which is what enables stateless per-pixel execution. The parameters are populated freshly each frame by `launch_filter`:

- **float alpha** blend intensity multiplier. In the CLI, this oscillates between 1.0 and 3.0 (0.01/frame). Used by blending and alpha-scaling filters to modulate mix ratios. At 1.0, blends are subtle; at 3.0, blends are oversaturated and aggressive.
- **bool isNegative** when true, the `setAlpha` device function inverts all three color channels (255 - value) after each pixel's filter chain completes. This globally toggles a photographic negative effect that stacks on top of any filter output.
- **int numFrames** number of valid entries in the history ring (up to `arraySize`). Filters that index into `allFrames[]` use this to avoid reading uninitialized slots before the buffer is full.
- **int square_size** block tile dimension in pixels, oscillates between 2 and 64. Controls the tile/block granularity for effects like `SquareBlockResize`, `DiagPixelated`, `pixelScale`, and others that work in rectangular pixel regions.
- **int start_index** the current bouncing frame selection index (0 to `arraySize-1`). Passed directly to effects that use it as a frame history lookup index, creating smooth temporal sweeping.
- **int start_dir** the direction of the frame index bounce (1 = forward, 0 = reverse). Some effects use this to modulate the direction of their temporal lookup.
- **int int_param1** `rand() % height`. Used as a random vertical offset for pixel shift effects (`PictureShiftDown`, `DistortedLinesY`, etc.). Changes every frame so shifts appear random and non-repeating.
- **int int_param2** `rand() % width`. Same concept as `int_param1` but horizontal. Used for horizontal shift and distortion effects.
- **float float_param1** `(float)(rand() % 255)`. A random brightness or gradient value. Used in effects like `GradientFlashColor` where a random scalar modulates the gradient output per frame.
- **int seed** `rand()`. Per-frame random seed passed to `gpu_rand(x, y, seed)` inside device functions. Since `gpu_rand` is deterministic given the same (x, y, seed), all threads with the same seed get the same noise map but the map changes every frame because the seed is regenerated.
- **int frame_count** monotonically increasing counter (`frame_counter++`). Used for time-based oscillations inside device functions (e.g., `processTripHSV` uses it to advance HSV hue over time, `get_osc_offset` uses it for sinusoidal offsets). This is the primary time variable for all internally animated kernels.
- **int sumR, sumG, sumB** each set to `rand() % 255` per frame. Used in bitwise XOR effects like `XorSumStrobe` where the pixel's channels are XOR'd against a random color triplet, creating a color-shifting strobe.

- **int threshold** fixed at 15. Used by `colorBounds()` in effects that do color comparison (e.g., `StrangeGlitch` switches to a historical frame pixel when the current and historical color differ by more than the threshold). Tuned to ignore slight color noise while catching meaningful motion.
- **int sw, sh** dynamic block dimensions: `16 + (frame_counter % 48)`. Creates a slowly growing/shrinking block size for diagonal pixelated effects like `DiagPixelatedResize`. The range is 16 to 63 pixels wide/tall, cycling over 48 frames.

The Unified Kernel `unifiedFilterKernel`

The `__global__ void unifiedFilterKernel` is the single CUDA kernel that executes the entire filter chain per frame. It is the pixel-level engine of the whole project. Understanding how it works explains why the system can run any combination of 736 effects in real-time.

Thread Layout and Pixel Assignment

Each CUDA thread is assigned one pixel:

- `int x = blockIdx.x * blockDim.x + threadIdx.x`
- `int y = blockIdx.y * blockDim.y + threadIdx.y`

With `blockDim = (16, 16)`, each block contains 256 threads arranged as a 1616 tile of pixels. A 19201080 frame requires `ceil(1920/16) * ceil(1080/16) = 120 * 68 = 8160` blocks. At 256 threads per block, that is 2,088,960 threads one per pixel all executing simultaneously on the GPU.

A bounds check (`if (x >= width || y >= height) return`) handles edge tiles where the grid doesn't align perfectly with frame dimensions.

The Filter Chain Loop

Inside each thread's pixel scope, the kernel loops over the filter list:

```
for (int i = 0; i < count; ++i) {
    switch (filters[i].index) {
        case 0: processSelfAlphaBlend(x, y, data, step, params); break;
        case 1: processMedianBlend(x, y, data, allFrames, step, params); break;
        // ... 734 more cases ...
        case 735: processHueWobble(x, y, data, step, params); break;
    }
}
```

Each `process*()` call is a `__device__` function that reads and writes the pixel at `data[y * step + x * 4]` in RGBA byte order. Because all threads within a warp execute the same filter index at the same time, the switch dispatch is warp-coherent when all threads in the warp are processing the same filter which they always are, since the filter list is shared params, not per-pixel data.

After the loop completes, `setAlpha(data, y * step + x * 4, params.isNegative)` is called once to finalize the alpha channel and optionally negate the color.

Device Helper Functions

- `__device__ float gpu_rand(int x, int y, int seed)` a deterministic per-pixel hash function. Uses three multiply-and-XOR operations seeded by x, y, and the per-frame seed, followed by XOR-shift scrambling, returning a float in [0, 1). Used by noise effects (`VisualSnow`, `StaticGlitch`, `RandomPixels`, `DistortedLines`) to generate pixel-level noise without any global shared state, which is important since CUDA threads cannot safely share mutable state.
- `__device__ void setAlpha(unsigned char* data, int idx, bool isNegative)` if `isNegative`, inverts R, G, B (three bytes at `idx`, `idx+1`, `idx+2`). Always writes 255 to the alpha byte (`idx+3`), ensuring fully opaque output regardless of what filters wrote to that byte.

- `__device__ bool colorBounds(r1,g1,b1, r2,g2,b2, ir,ig,ib)` returns true if the absolute difference of each channel pair is within the given threshold. Used by motion-detection effects: **StrangeGlitch** uses it to detect when a pixel has changed significantly between frames and swap in the historical pixel value; **MatrixOutline** uses it to zero out pixels that match a reference frame (creating a motion-outline effect).

The Filter Catalog 736 Effects Across All Categories

The filter table is defined at the top of `filters.cu` as a static array of `Filter` structs. Index 0 through 735 are defined. Each filter has its own `__device__` implementation function but all are dispatched through the single unified kernel switch. Here is a breakdown of the major effect families that make up the library:

Temporal / Trail Effects

Read from historical frames in the `DynamicFrameBuffer` ring to create motion persistence, ghosting, and smear effects.

- **MedianBlend** averages current pixel against all history frames + XOR + contrast boost
- **AuraTrails** blends against frames at history indices 1, 4, and 7
- **MotionGhostTrails** linear blend between current and oldest frame via alpha
- **WaveTrails** wave-modulated temporal blend
- **RGBLineTrails** / **RGBWideTrails** / **RGBLongTrails** per-channel trail variations with different decay widths
- **ProperTrails** / **ShortTrail** clean blend-from-history patterns
- **AcidTrailsBlend** / **GhostTrailsBlend** higher-level composites of trail motion
- **EchoBlend** / **EchoShift** / **TrailEcho** echo-delay style history reads
- **FrameBlendMulti** / **FrameBlendMultiX** blending across multiple history slots simultaneously

Geometric Distortion

Remap pixel coordinates shift, stretch, warp, flip, tear, and distort spatial layout.

- **PictureShiftDown** / **PictureShiftRight** / **PictureShiftVariable** whole-frame translate by random pixel offset
- **StretchR/G/B_Right** / **StretchR/G/B_Down** per-channel horizontal or vertical stretch
- **TearRight** / **TearDown** VHS-style tape tear distortions
- **DistortionByRow** / **DistortionByCol** row/column-level random displacement
- **RippleEffect** / **ShockWave** / **TwistEffect** radial math-based spatial warps
- **FishEye** / **TunnelEffect** / **VortexEffect** / **Kaleidoscope** polar coordinate remap effects
- **SpiralWave** / **SpiralTrail** spiral coordinate mapping
- **ZoomBlur** / **RadialBlur** / **RotateBlend** zoom and rotational blur patterns
- **ExpandContract** / **ExpandLeftRight** / **DiagInward** pixel expansion/contraction patterns
- **MirrorWave** / **MirrorWaveX** / **MirrorWaveY** wave-modulated mirror remap

Color Manipulation

Alter pixel color values mathematically without spatial remapping.

- **SelfAlphaBlend** multiplies each channel by (1 + alpha)
- **SelfScaleRefined** / **SelfScaleByFrame** channel scaling with clamping
- **TripHSV** converts to HSV, cycles hue by frame_count, converts back
- **GradientRainbow** / **GradientSelf** / **GradientDown** / **GraidentHorizontal** positional gradient overlays
- **FadeRtoGtoB** / **FadeRGB_Speed** / **FadeRandomChannel** sequential channel fade patterns
- **ColorAccumulate1/2/3** / **colorAccumulate** accumulation blend series
- **HueRotate** / **ChromaticAberration** / **RGBShift** color space shift effects

- `Posterize` / `Solarize` / `GammaBright` / `GammaDark` / `ContrastBoost` / `ContrastReduce` tone-mapping operations
- `TruncateColor` / `TruncateVariable` / `TruncateVariableScale` color depth reduction patterns
- `ColorDrift` / `ColorPulse` / `ColorPulseRGB` / `ColorFadeFilter` time-animated color drift series

Bitwise / XOR Operations

Apply bitwise logic (XOR, AND, OR) between pixel channels, historical frames, or random values.

- `Bitwise_XOR` / `Bitwise_AND` / `Bitwise_OR` current frame XOR/AND/OR with history frame
- `XorSumStrobe` XORs each channel against a random sumR/G/B value
- `XorAlpha` / `XorFade` / `XorSine` / `XorLag` / `XorScale` XOR variant series
- `SelfXorBlend` / `SelfXorDoubleFlash` / `SelfOrDoubleFlash` self-XOR patterns
- `BitwiseXorStrobe` / `BitwiseRotateBlend` / `BitwiseXorScaleBlend` scaled and rotated XOR blends
- `AndStrobe` / `AndStrobeScale` / `AndOrXorStrobe` / `AndOrXorStrobeScale` AND-based strobe family
- `MedianBlendXor` / `CollectionAlphaXor` XOR composited with median blend

Glitch and Noise

Simulate digital corruption, static, tracking errors, and signal noise.

- `StrangeGlitch` color-bounds detection, swaps pixel with history when change detected
- `HorizontalGlitch` / `VerticalGlitch` glitch lines by row or column history sampling
- `PixelGlitch` / `StaticGlitch` / `LineGlitch` / `BoxGlitch` progressive glitch types
- `GlitchBlock` / `GlitchBlockXor` / `GlitchLine` / `GlitchLineX` block and line glitch family
- `VisualSnow` / `VisualSnowX2` `gpu_rand`-based noise overlay
- `NoiseBlend` / `NoiseBlendX2` / `NoiseXor` / `NoiseBlendX` noise blend family
- `VHSTracking` emulates VHS head-tracking horizontal tear artifact
- `DataCorrupt` / `DigitalArtifact` / `TapeGlitch` / `ColorGlitch` digital corruption effects
- `SliceGlitch` / `GlitchSort` / `GlitchMosaic` sorted and mosaic-style glitch

Scan / Line Effects

Work on scanline-level patterns, interlacing, and whole-row/column operations.

- `InvertedScanlines` / `ScanSwitch` / `ScanAlphaSwitch` scanline inversion and toggle
- `HorizontalLines` / `DiagonalLines` / `BlackLines` / `LongLines` line overlay patterns
- `LineInLineOut` series (3947, 262264) complex oscillating line-in/line-out passes
- `YLineDown` / `YLineDownBlend` vertical line drift effects
- `LineGlitch` / `LineGlitchX` / `LineAcrossX` / `LinesAcrossX` line-based glitch patterns
- `BlendedScanLines` / `InterlaceBlend` interlace-style line blending
- `ShiftLinesDown` whole-scanline vertical displacement

Block / Square Effects

Operate on rectangular pixel regions pixelation, block swapping, tile-based transforms.

- `SquareBlockResize` divides frame into horizontal bands, each blended against a different history frame
- `SquareShrink` clamps inner region against history frame based on oscillating offset
- `SquareByRow` / `SquareByRowRev` / `SquareByRow2` / `SquareByRow2Plus` row-ordered block history sampling
- `DiagPixelated` / `DiagPixelatedResize` diagonal block averaging with fixed or variable tile size
- `PixelateBlend` / `PixelateRect` / `PixelateWave` / `MosaicBlend` pixelation and mosaic blends
- `BlockPixels` / `BlockScale` / `BlockXor` / `BlockStrobe` block-level transform series
- `BlockyTrails16` / `BlockyTrails32` block-sized history trails

- **BlockSwap** swaps pixel blocks between current and history

Wave / Oscillation Effects

Use sine, cosine, square, triangle, sawtooth, and pulse waveforms to drive pixel displacement or color modulation.

- **SineWaveDistort** / **CosineWaveDistort** horizontal/vertical sinusoidal pixel shifting
- **SineWaveBlend** / **CosineWaveBlend** / **SinCosBlend** wave-modulated blends
- **SquareWave** / **SquareWaveX** / **SquareWaveBlend** square wave color/blend modulation
- **TriangleWave** / **TriangleWaveBlend** triangle wave modulation
- **SawtoothWave** / **SawtoothWaveBlend** sawtooth wave modulation
- **PulseWave** / **PulseWaveBlend** / **PulseRadial** / **Pulse** pulse-shape modulation
- **StepWave** / **StepWaveBlend** step-function wave modulation
- **WaveBlend** / **WaveBlendX2** / **WavePattern** / **WavePatternXor** generic wave blend series
- **MirrorWave** / **SpiralWave** / **VortexEffect** / **TwistEffect** complex wave-geometry hybrids

Mirror / Reflection

Symmetric transforms that fold, flip, or reflect pixel coordinates.

- **MirrorReverseColor** four-point mirror average (top-left, bottom-left, bottom-right, current) with channel reversal
- **AlphaBlendMirror** / **MirrorXorAlpha** mirror with alpha and XOR compositing
- **IntertwinedMirror** interleaved mirror blend
- **MirrorReverseColorBlend** blended version of **MirrorReverseColor**
- **FlipAlphaBlend** / **RandomFlipFilter** / **FlipPictureShift** / **FlipMirror** flip transform series
- **DiagMirror** / **ShadowMirror** / **GhostMirror** / **FacetMirror** diagonal and ghost mirror types
- **Kaleidoscope** / **KaleidoscopeBlend** / **KaleidoBlend** / **KaleidoScope4** / **MirrorKaleid** kaleidoscope family
- **SplitMirror** / **TripleSplit** / **PrismSplit** split-panel mirror effects

Pixel Read / Strobe / Blend Collections

Effects that do frame-history collection reads, random collection sampling, or strobe-pattern switching.

- **StretchColMatrix8/16/32** samples history at (x / sw) % numFrames column stride
- **ColorCollectionSubtleStrobe** / **CollectionRandom** / **CollectionAlphaXor** collection random/strobe series
- **ColorCollection64X** / **ColorCollectionSwitch** / **ColorCollectionRGB_Index** indexed collection blends
- **ColorCollectionGhostTrails** / **ColorCollectionScale** / **ColorCollectionReverseStrobe** ghost and scale variants
- **BlendWithSource25/50/75/100** fixed-ratio blends with history source
- **BlendFor360** / **BlendForward16** / **BlendForward32** / **BlendFromXtoY** direction-controlled blend series
- **FadeInAndOut** / **FadeBlendXor** / **FadeBars** fade pattern series
- **MildStrobe** / **BrightStrobe** / **DarkStrobe** / **StrobeEffect** / **StrobeXor** strobe type family

Advanced / Cinematic Effects

Higher-complexity effects that combine multiple techniques or simulate specific visual phenomena.

- **CRTCurvature** simulates CRT screen barrel distortion
- **FilmGrain** per-pixel noise scaled by luma to simulate analog grain
- **ChromaticAberration** / **ChromaticAberrationX** RGB channel lateral displacement simulating lens chromatic error

- `LensFlare` / `LightLeak` / `GlowTrails` / `NeonGlow` / `GlowEdge` / `GlowPulse` optical glow/flare family
- `NightVision` / `InfraredView` / `ThermalBlend` synthetic imaging modalities
- `MatrixCode` / `DigitalRain` column-drip pattern simulating cascading character streams
- `WaterColor` / `OilSlick` / `LiquidMetal` / `LavaLamp` fluid/material simulation-inspired effects
- `HeatDistort` / `HeatRipple` / `HeatWave` heat shimmer displacement series
- `SobelNorm` / `DetectEdges` / `SketchOutline` / `SobelGlow` / `ElectricEdge` edge detection and outline family
- `GalaxySpiral` / `CosmicDust` / `StarBurst` / `Fireworks` space/particle aesthetic effects

The Two-Stage GPU Pipeline CUDA Compute + GLSL Shaders

The systems visual power comes from combining two distinct GPU programming models into a single per-frame pipeline. Understanding how each stage works and how they connect explains why the project can produce such a large array of effects.

Stage 1: CUDA Compute Kernels Massively Parallel Pixel Processing

NVIDIA GPUs contain thousands of streaming multiprocessors (SMs), each capable of running many threads in parallel. Both CUDA kernels and GLSL shaders execute on this same hardware the difference is the programming model. `acidcam-gpu` uses the **CUDA compute model** to assign **one thread to each pixel** of the video frame. For a 1920x1080 frame, the system launches **2,088,960 threads simultaneously**, organized into 1616 blocks of 256 threads each (8,160 blocks total). Every pixel in the frame is processed in parallel there is no sequential per-pixel loop on the CPU.

Each thread runs the **unified dispatch kernel** (`unifiedFilterKernel`), which loops through the users selected filter list and applies each filter to that threads pixel in order. The kernel contains a 736-case `switch` statement each case calls a `__device__` function that reads and writes the pixels RGBA values at `data[y * step + x * 4]`. Because every thread in a warp processes the same filter index at the same time (the filter list is shared, not per-pixel), the dispatch is **warp-coherent** and runs at full GPU throughput.

Crucially, the CUDA filters have access to a **ring buffer of historical frames** stored entirely in GPU device memory (`DynamicFrameBuffer`). This means filters like `MedianBlend` can average across all prior frames, `AuraTrails` can blend against frames at specific history indices, and `MatrixOutline` can compare against a frame from several steps back all without any data leaving the GPU. This kind of arbitrary random-access memory read across multiple frame buffers is what makes the CUDA compute model essential GLSL fragment shaders run on the same GPU cores but operate within the graphics pipeline, which doesnt support this kind of free-form device memory access.

Stage 2: GLSL Shaders Full-Frame Post-Processing

After CUDA finishes processing every pixel, the result must be displayed. Rather than downloading the frame to the CPU and re-uploading it, the system uses **CUDA/OpenGL interop via a Pixel Buffer Object (PBO)**:

1. The PBO is registered with CUDA using `cudaGraphicsGLRegisterBuffer`
2. CUDA maps the PBO into its address space and performs a `cudaMemcpy2D` (device-to-device) from the working buffer into the PBO **the frame never touches CPU RAM**
3. OpenGL binds the PBO as a pixel unpack buffer and calls `glTexSubImage2D` to populate a texture another zero-copy GPU operation

The CUDA-filtered frame is now an OpenGL texture. At this point, **GLSL fragment shaders** take over. These are standard OpenGL shader programs (GLSL 330 core) that process the entire frame as a texture. The shader library provides Shadertoy-compatible uniforms `iTime`, `iResolution`, `iMouse`, `iFrame`, `iTimeDelta`, `iDate`, `iFrameRate`, `iSampleRate`, plus audio-reactive uniforms `amp/uamp` enabling a wide range of full-frame effects like color grading, distortion warps, glow effects, CRT simulation, and procedural pattern overlays.

Shaders can render the texture onto a **fullscreen 2D quad** (standard image processing) or onto **3D model geometry** loaded from MXMOD files, projecting the filtered video onto a rotating or animated mesh.

The Transfer: Zero-Copy GPU Interop

The bridge between CUDA and OpenGL is critical to performance. The `TextureUploader` class handles this via PBO interop the CUDA working buffer is copied directly into an OpenGL PBO in device memory, then uploaded to a GL texture, all without a CPU round-trip. This means the entire pipeline from raw camera frame to final displayed output can stay on the GPU throughout both the CUDA and GLSL stages.

Filter and Shader Stacking Two Independent Composable Chains

The system provides **two independent stacking mechanisms** that compose together: a CUDA filter chain and a GLSL shader pass stack. Users can configure both, and every combination produces a different visual result.

CUDA Filter Chain (Stacking Inside the Kernel)

Users select any subset of the 736 available CUDA filters and arrange them in a specific order. This ordered list is uploaded to GPU memory as an array of `GPUFilter` structs. Inside the kernel, every pixel thread loops through this array:

```
for (int i = 0; i < count; ++i) {
    switch (filters[i].index) {
        case 0: processSelfAlphaBlend(x, y, data, step, params); break;
        case 1: processMedianBlend(x, y, data, allFrames, step, params); break;
        // ... 734 more cases ...
    }
}
```

Each filter modifies the pixel buffer **in-place**. Filter N reads the output that filter N-1 wrote. This means the filters chain naturally the output of one becomes the input of the next, all within a single kernel launch. There are no intermediate buffer copies between filters; the stacking is purely sequential modification of the same pixel data.

Rebuilding the filter list is **lazy and zero-cost at steady state**: the device-side filter array is only reallocated when the chain actually changes (tracked by a **changed** flag). Re-ordering filters at runtime is essentially free.

GLSL Shader Pass Stack (Ping-Pong Framebuffer Passes)

After CUDA processing, GLSL shaders can be stacked using a **multi-pass ping-pong framebuffer** technique. The user selects and orders GLSL shader passes through the `ShaderPassDialog` interface. Each pass works as follows:

1. The system maintains two offscreen framebuffers (`passFB0[0]` and `passFB0[1]`) with corresponding textures (`passTexture[0]` and `passTexture[1]`)
2. The input starts as the camera texture (which already contains the CUDA-filtered output)
3. For each shader in the pass list: - Bind `passFB0[pingpong]` as the render target - Activate the pass shader program and set its uniforms (time, resolution, etc.) - Draw a fullscreen quad sampling from `inputTex` - Set `inputTex = passTexture[pingpong]` (the output becomes the next pass input) - Flip `pingpong = 1 - pingpong`
4. The final texture becomes the input for the main render (2D sprite or 3D mesh)

This allows stacking any number of GLSL shader effects each one processes the output of the previous, building up complexity layer by layer.

Combined Stacking: The Full Pipeline

Both chains compose into a single per-frame pipeline:

Camera Frame

```
[CUDA Filter 1] [CUDA Filter 2] ... [CUDA Filter N]    (in-kernel chain)
PBO interop (zero-copy GPU transfer)
[GLSL Shader Pass 1] [GLSL Shader Pass 2] ... [GLSL Pass M]    (ping-pong FBOs)
Final composite (2D quad or 3D mesh)
Screen / Recording
```

The user independently controls both the CUDA filter list and the GLSL shader pass list. Changing either one or just reordering elements produces a completely different visual result.

Why Stacking Produces Infinite Visuals

Each filter and shader is a transformation function. Ordering is not commutative in general, so:

```
F3(F2(F1(frame))) != F1(F2(F3(frame)))
```

If you choose n distinct filters and order them, permutations are $n!$. Even before parameter changes, feedback buffers, and time-varying uniforms, this grows explosively.

- 5 filters 120 orderings
- 8 filters 40,320 orderings
- 10 filters 3,628,800 orderings

With 736 CUDA filters to choose from, plus an independent library of GLSL shaders that also stack and reorder, the combinatorial space multiplies further. Add in the continuously evolving per-frame parameters (alpha oscillation, frame history index, square size, random seeds, frame count) and the visual output becomes effectively unbounded the same filter chain produces different results every frame because the parameters driving it are always changing.

In-Depth: Why Order Dominates Output

In `unifiedFilterKernel`, each thread applies filter case statements in strict list order. Many cases are stateful across time because they read `allFrames` (historical frame pointers) and per-frame randomized params generated in `launch_filter`.

- **Non-commutative transforms:** geometric shift then color XOR yields different values than color XOR then geometric shift, because pixel neighborhoods sampled differ.
- **Temporal dependency:** trail/motion filters sample prior frames; changing order changes which transformed history is fed forward.
- **Parameter evolution:** alpha, square size, random seed, and thresholds evolve frame-to-frame, so visual output is a time series, not a static map.
- **Resolution coupling:** many filters compute indices with width/height and step; same chain on different resolutions creates different artifact structure.
- **Cross-domain stacking:** CUDA filter chains and GLSL shader pass stacks compose independently, so the total combinatorial growth is the product of both far larger than either alone. A 5-filter CUDA chain a 3-shader GLSL stack already produces $120 \cdot 6 = 720$ distinct orderings before any parameter variation.

Build System, Dependencies, and Distribution

The project requires a very specific dependency stack not all of it is available pre-built from most package managers. OpenCV must be compiled from source with CUDA support enabled, which is the most significant build-time requirement.

Required Dependencies

- **NVIDIA GPU:** RTX 20-series or newer. The project is developed and optimized on an RTX 2070. Older hardware may work but is not guaranteed.
- **NVIDIA Proprietary Drivers:** v535 or newer. The CUDA runtime links against driver libraries at version-specific interfaces.
- **CUDA Toolkit 12.x:** provides `nvcc`, `cuda_runtime.h`, `cudaMalloc/cudaMemcpy/cudaDeviceSynchronize`, and device code compilation.
- **OpenCV with CUDA support:** must be compiled from source with `WITH_CUDA=ON` and matching CUDA architecture flags. Provides `cv::cuda::GpuMat`, `cv::cuda::cvtColor`, `cv::cuda::printShortCudaDeviceInfo`, and the `cv::VideoCapture` backend used in the CLI.
- **libmx2 / MXWrite:** the MX2 ecosystem library provides MXMOD 3D model parsing and the MXWrite encoder wrapper around FFmpeg. Must be built and installed before building `acidcam-gpu`.
- **Qt6:** required only for the ACMX2 Qt interface, not for the library or CLI.
- **C++20 compiler:** the project uses `std::format`, `std::filesystem`, and other C++20 features throughout `main_cv.cu`.

Build Order

1. Build and install `libmx2/libmx` (with OpenGL support): `cmake .. -DEXAMPLES=OFF -DOPENGL=ON && make -j$(nproc) && sudo make install`
2. Build and install `acidcam-gpu/MXWrite`: `cmake .. && make -j$(nproc) && sudo make install`
3. Build and install the `acidcam-gpu` library and CLI: `cmake .. && make -j$(nproc) && sudo make install`. This installs the shared library, headers, and CMake package config so downstream projects can use `find_package(acidcam-gpu)`.
4. Optionally build and install the Qt ACMX2 interface.

Container Distribution (Podman)

The primary distribution mechanism for end users is a Podman container image: `ghcr.io/lostjared/acmx2:latest`. This avoids requiring users to compile OpenCV from source. The container includes all runtime dependencies including NVIDIA CUDA libraries (subject to NVIDIA's Deep Learning Container License). To run it:

- Host must have NVIDIA drivers and NVIDIA Container Toolkit for Podman installed.
- The run script (`podman/run-acmx2.sh`) passes GPU access, camera device (`/dev/video0`), audio device, and X11 display socket into the container so the app appears on the host desktop with full hardware access.
- This approach means users on Bazzite, Arch, or other NVIDIA-equipped Linux systems can run a fully GPU-accelerated visual effects tool without any build steps.

Development Environment

The project is developed on **Bazzite Linux** using **Arch Linux containers via Distrobox**. This means the host OS is an immutable Fedora-based image, and all development tooling (CUDA, OpenCV from AUR, Qt6, gcc) is installed inside an Arch container mounted into the same home directory. This pattern keeps the base OS clean while maintaining full package manager access for development dependencies.

Source Code Map by Project Part

ACMX2 Core

- `ACMX2/CMakeLists.txt` dependency checks + executable wiring
- `ACMX2/acmx.cpp` capture, render, controls, runtime loop
- `ACMX2/program.cpp/.hpp` shader program loading + binary cache
- `ACMX2/data/*.glsl` base passthrough/framebuffer shaders

ACMX2 Media + Audio + 3D

- `ACMX2/audio.cpp/.hpp` RtAudio amplitude/reactive hooks
- `ACMX2/audio_transfer.cpp` audio transfer helper
- `ACMX2/models/*.mxmod` MXMOD geometry library
- `ACMX2/examples/*.glsl` sample shader library/index

ACMX2 Interface + Tools

- `ACMX2/interface/*` Qt launcher/editor/dialog workflow
- `ACMX2/shader_generator/*` AI-assisted shader generation utility
- `ACMX2/shader.packs/` shader pack metadata
- `ACMX2/macos/` platform-specific notes

acidcam-gpu Library

- `acidcam-gpu/CMakeLists.txt` CUDA/OpenCV/MXWrite package build
- `acidcam-gpu/include/ac-gpu/ac-gpu.hpp` API structs + launch contract
- `acidcam-gpu/src/filters.cu` filter table + unifiedFilterKernel dispatch
- `acidcam-gpu/app/main_cv.cu` CLI runtime/argument pipeline

Ops + Deployment

- `podman/Containerfile` containerized runtime image
- `podman/run-acmx2.sh` GPU/camera/audio passthrough run script
- `acidcam-gpu/scripts/*` OpenCV CUDA and environment helper scripts

Expanded Meaning of Each Code Map Item

- **ACMX2/CMakeLists.txt:** hard-gates feature availability at configure time; this file decides whether your build can legally run all advertised runtime paths.
- **ACMX2/acmx.cpp:** the largest operational core; input parsing, frame lifecycle, cache management, GL/CUDA interop, and output/control behavior converge here.
- **ACMX2/program.cpp/.hpp:** shader binary cache layer uses source+driver fingerprinting to skip repeated compile/link cycles and reduce startup latency.
- **ACMX2/data/*.glsl:** baseline vertex/fragment units that form base render stages and buffer transfer steps.
- **ACMX2/audio.cpp/.hpp:** real-time amplitude extraction from audio input; callback computes average magnitude used for reactive controls.
- **ACMX2/audio_transfer.cpp:** utility path for audio transfer/record style workflows and supporting media synchronization tasks.
- **ACMX2/models/*.mxmod:** model assets for 3D or scene-influenced visuals integrated with runtime rendering.
- **ACMX2/examples/*.glsl:** curated shader samples used as practical templates and quick-start visual blocks.
- **ACMX2/interface/*:** user-facing orchestration UI for session setup, process launch, list reordering, and settings persistence.
- **ACMX2/shader_generator/*:** assistant tooling that helps generate shader ideas/workflows while maintaining runtime-compatible output.
- **ACMX2/shader.packs/:** pack metadata and organization layer for large shader collections.
- **acidcam-gpu/include/ac-gpu/ac-gpu.hpp:** ABI/API boundary between host app and CUDA engine.
- **acidcam-gpu/src/filters.cu:** massive effect implementation + dispatch switch; this file is the core visual transformation engine.
- **acidcam-gpu/app/main_cv.cu:** standalone CLI proving the CUDA pipeline without the full ACMX2 UI stack.

- **podman/Containerfile and scripts:** reproducible deployment path for camera/GPU-enabled container runs.

End-to-End Runtime Flow

1. Launch ACMX2 or the Qt interface.
2. Load camera/video + shader set + optional GPU filter list.
3. Frames enter OpenCV/CUDA buffers.
4. Ordered filter list is sent into CUDA dispatch.
5. GPU writes transformed pixels back for GL presentation.
6. Optional recording/snapshot/audio-reactive controls are applied.

The key design principle is: **order is the language**. Re-ordering the same operators creates a different visual grammar every time.

Detailed Runtime Inner Workings

1. **Configuration phase:** CLI args or Qt controls define input source, CUDA filters, shader pass order, buffer depth, fps/recording, and optional audio mode.
2. **Device/stream validation:** runtime checks CUDA device count, camera/video availability, and opens required streams; invalid resources fail early with explicit messages.
3. **Frame ingress:** each frame enters as OpenCV matrix data; format conversions normalize channel layout for downstream CUDA/OpenGL pipelines.
4. **History update:** dynamic or fixed frame buffers rotate so current + historical frames are available in device memory simultaneously.
5. **Parameter evolution:** animated values (alpha ramps, square size oscillation, frame index direction) are updated once per frame tick.
6. **Filter dispatch packaging:** selected filter IDs are transformed into GPU list structs; if chain changed, device filter list is rebuilt.
7. **CUDA execution:** unified kernel runs one thread per pixel, then loops each selected filter in order, modifying pixel values cumulatively.
8. **Interop transfer:** resulting CUDA buffer is copied into GL-bound PBO memory and uploaded to texture for immediate rendering without host bounce.
9. **Output fan-out:** displayed frame may also be written to file and/or used by snapshot threads; optional audio state can modulate rendering behavior.
10. **Loop continuation:** next frame repeats with updated temporal state, producing evolving visuals that are path-dependent over time.

Project Architecture Goals and Design Tradeoffs

The project balances three competing goals: maximal visual complexity, real-time responsiveness, and a workflow that remains controllable by artists/operators. The architecture is intentionally modular so each subsystem can evolve without collapsing the full stack.

1) Composability First

Shaders, CUDA filters, and temporal buffers are treated as composable operators. This favors creative exploration and makes the system useful for live experimentation, not just fixed presets.

2) Deterministic Runtime Paths

The same ordered chain and parameter set produce reproducible behavior at a given resolution/device configuration, which matters for iterative artistic workflows and debugging.

3) Throughput-Oriented Execution

CUDA/OpenGL interop and device-side frame history reduce host-device copies. The core design prefers sustained frame throughput over heavyweight per-frame orchestration.

4) Operator Visibility

The Qt layer exposes logs, ordered lists, and persistent settings. This turns complex GPU behavior into a controlled session system that can be repeated and tuned.

Engineering Implications

- **Memory vs flexibility:** temporal filters require multiple historical frames in VRAM; this increases memory usage but enables richer motion feedback effects.
- **Unified kernel benefits:** one dispatch model simplifies orchestration and lets filter chains be data-driven rather than compile-time fixed.
- **Cross-platform pressure:** CMake + container scripts improve portability, but GPU/audio/camera stacks still vary by OS and driver quality.
- **UI/runtime separation:** Qt remains a control plane while CLI/CUDA remains an execution plane; this separation keeps debugging cleaner.
- **Creative-state persistence:** cached settings and ordered chains make sessions restorable, which is important for long-form visual composition.

Build, Deployment, and Operational Model

Beyond visuals, the project includes practical operational pathways for repeatable builds and runtime deployment across development workstations and containerized environments.

Build Surface

- **Dependency validation:** CMake scripts act as a gatekeeper so unsupported environments fail during configure/build, not deep in runtime.
- **Library/executable layering:** the CUDA library can be tested through its own CLI path, while ACMX2 and Qt consume it as a higher-level orchestration stack.
- **Shader and model assets:** runtime behavior depends not only on binaries but also on curated GLSL and MXMOD assets shipped with the project.

Operational Surface

- **Container workflows:** Podman files and scripts define a reproducible execution recipe including GPU/camera/audio passthrough where supported.
- **Session-driven usage:** operators can treat configurations as sessions, allowing repeatable live setups for streaming, VJ work, or iterative rendering passes.
- **Failure visibility:** logs from runtime and UI layers help pinpoint whether issues are input-device, shader compile, CUDA dispatch, or encoder related.

Project Summary

acidcam-gpu + ACMX2 is a layered real-time visual computing system built entirely around NVIDIA CUDA. At the bottom: a 9,165-line CUDA file (`filters.cu`) implementing 736 per-pixel GPU effects dispatched through a single unified kernel. In the middle: a rotating device-side frame history buffer (`DynamicFrameBuffer`), a per-frame evolving parameter bundle (`FilterParams`), and a lazy-rebuild filter list mechanism that makes chain reconfiguration zero-cost. At the top: a CLI application with an oscillating `AnimationState` that continuously evolves alpha, block size, and temporal frame selection, plus a Qt6 orchestration interface with session persistence, process supervision, and multi-pass shader support. Distribution via Podman containers removes the painful dependency build requirement for end users.

The core creative insight is that **filter order is the language**. The same 736 effects, ordered differently, with different animated parameters, on different input material, produce an effectively unbounded space of visual outcomes all running at real-time framerates on a single consumer NVIDIA GPU.